

# Learning to Extract Coherent Summary via Deep Reinforcement Learning

Anonymous authors

## Abstract

Coherence plays a critical role in producing high-quality summary from a document. In recent years, neural extractive summarization is becoming increasingly attractive. However, most of them ignore the coherence of summary when extracting sentences. As an effort towards extracting coherent summaries, we propose a neural coherence model to capture the cross-sentence semantic and syntactic coherence patterns. The proposed neural coherence model obviated the need for feature engineering and can be trained from scratch in an end-to-end fashion using a large corpus. Empirical results show that the proposed neural coherence model can effectively capture the cross-sentence coherence patterns. Using the output of the neural coherence model and ROUGE as the reward, we design a reinforcement learning method to train the neural extractive summarizer which is named Reinforced Neural Extractive Summarization (RNES) model. The RNES learns to optimize coherence and informative importance of the summary simultaneously. Experiment results show that the proposed RNES outperforms existing baselines and achieves the state-of-art performance in terms of ROUGE on CNN/Daily Mail dataset. The qualitative evaluation indicates that the summary produced by the proposed RNES is more coherent and readable.

## Introduction

Although deep neural networks have dominated almost every field of natural language processing, such as sentiment classification (Tang, Qin, and Liu 2015), machine translation (Bahdanau, Cho, and Bengio 2014), generating high-quality summaries from long documents is still a very challenging task. Most of the recent works on abstractive summarization focus on headline generation from one paragraph (Rush, Chopra, and Weston 2015) or several sentences (Hu, Chen, and Zhu 2015) by using sequence-to-sequence architectures borrowed from neural machine translation. However, all of them bypass the fundamental problems of abstractive summarization, namely the representation of long documents and consecutive sentences generation. These models can not generate readable, informative and coherent sentences when dealing with full documents. There is still a long way for abstractive summarization to be practicable.

In contrast, extracting sentences from document to form summary, also named extractive summarization, is a more practical way, because it can guarantee the grammatical correctness and semantic relevance with the document of the produced summary. Extractive summarization has been researched for several decades. Traditional approaches mainly focus on scoring sentences of the document by using a graph-based method (Erkan and Radev 2004) or integer linear programming (Berg-Kirkpatrick, Gillick, and Klein 2011), coupling with human-engineered features. As the distributed representation shows its outstanding performance to capture the semantic and syntactic information of text (Mikolov et al. 2013; Hu et al. 2016), there is an emergence of works that use the deep neural network to extract salient sentences (Nallapati, Zhai, and Zhou 2017; Nayeem and Chali 2017; Cheng and Lapata 2016). Although both of traditional methods and deep neural network based methods can identify the important sentences of the documents, both of them lead to the problem in overall coherence of the summary. Sometimes there are several topics in the summary and the interpretation of one sentence do not depend on its neighbors, which causes that the summary can not express a meaning as a whole.

The coherence of summary is essential for its readability and clarity. In the past decades, research works mainly focus on topical coherence. The most popular method is the entity grid model (Barzilay and Lapata 2008) which constructs a grid to represent grammatical and semantic transitions of entities between sentences. The entity grid is then converted to discrete feature vectors which are used as input to learning model to obtain the coherence score of the texts (Nguyen and Joty 2017). The entity grid model needs the entity labels of words in sentences. The performance of name entity recognition algorithm may become the bottleneck of the grid model. The entity grid model uses a discrete representation of text, which can not model long entity transitions. In this paper, we propose neural coherence model which can learn the coherence degree of two sentences with the distributed representation of text in an end-to-end fashion.

Due to the avoidance of lots of handcrafted features, the neural extractive summarization is becoming increasingly popular. The core contribution of most of these works is using the deep neural network to obtain the distributed representation of documents and sentences, to extract infor-

matively important sentences (Cheng and Lapata 2016). To our knowledge, there is no work incorporating the coherence into the neural extractive model while extracting sentences. It is difficult to incorporate coherence into optimization objective of supervised learning because coherence also depends on sentences that are eventually extracted when inference is performed. Reinforcement learning (RL) aims to train an agent to obtain the maximum reward when it interacts with a given environment. RL is usually used in the case that the evaluation metrics to be optimized is not differentiable and traditional supervised learning methods cannot be used. RL has been used in sequence generation tasks such as abstractive summarization generation (Paulus, Xiong, and Socher 2017) and neural machine translation by optimizing the metrics such as BLEU, ROUGE or METEOR (Nguyen, Boyd-Graber, and Daume 2017).

In this paper, we focus on improving the coherence ability of neural extractive model via reinforcement learning. Our first contribution is to propose a novel neural coherence model which uses the low-dimension dense distributed representation of texts instead of sparse discrete features. The proposed neural coherence model does not rely on any entity recognition systems and can be trained almost from scratch in an end-to-end fashion. The neural coherence model can capture the cross-sentence local entity transition and discourse relation with the layer by layer convolution and max-pooling layers. It is feasible to train the neural coherence model with large scale unlabeled plain text. The experiment results show that, given one sentence, the neural coherence model can identify the appropriate next sentence to compose a coherent sentence pair effectively.

Our second contribution is the reinforcement learning solution for training neural extractive summarization model. Although reinforcement learning has been applied to abstractive headline generation with Rouge or BLEU metrics as rewards (Paulus, Xiong, and Socher 2017; Ayana et al. 2016), to our knowledge, we are the first to design the Reinforced Neural Extractive Summarization (RNES) model with the neural coherence and rouge score as the reward. The proposed RNES model can balance the coherence and informative importance of sentences effectively. We experimentally evaluate our RNES model on standard task and show that it can achieve state of the art performance on rouge metrics. The qualitative evaluation indicates that the summary produced by RNES are more informative and coherent.

## Related Work

Our research is built on previous works in the field of neural extractive summarization, reinforcement learning and coherence models.

Much progress has been made beyond traditional frameworks of extractive summarization models. Most of recent works are based on deep neural networks and distributed representation of text, which regards the extractive summarization as sentence sequence labeling problem. For example, (Cheng and Lapata 2016) use convolutional neural network to encode sentence, and then recurrent neural network takes in the sentence representation sequentially to encode the document. Finally they use another recurrent neural

network to label sentences sequentially, taking the encoded document representation and the previously labeled sentences into account. (Cheng and Lapata 2016) mainly consider the importance of the sentence and the non-redundancy of the summary. (Nallapati, Zhai, and Zhou 2017) use the similar architecture to encode document, but add the information content, salience and novelty into the model when labeling sentences. The model proposed by (Nallapati, Zhai, and Zhou 2017) is very interpretable, since it allows visualization of its predictions.

Our work is also related to the application of reinforcement learning on document summarization. Different from text classification whose target is the clear single or multiple labels, the targets of machine translation or automatic text summarization are many potential sequence of words, which own its property independent of input. Such tasks are natural candidate problems for reinforcement learning by using some evaluation metrics to judge the quality of the result such as BLEU, ROUGE and et al. Deep reinforcement learning has drawn great attention in recent years. For example, (Paulus, Xiong, and Socher 2017) use the ROUGE-L score as a reinforcement reward and self-critical policy gradient training algorithm to train an abstractive summarization models. (Ayana et al. 2016) and (Ranzato et al. 2015) show optimizing the evaluation metrics ROUGE directly via reinforcement learning is more effective than optimizing likelihood for the decoder generation. All of above works focus on abstractive summarization by optimizing ROUGE metrics. To date, there are no work to apply RL with coherence as reward to neural extractive summarization. To our knowledge, our work is the first step toward this goal.

A key requirement for summarization systems is the coherence of its output. Coherence is which makes multiple sentences semantically, logically and syntactically coherent (Yao, Wan, and Xiao 2017). Not surprisingly, a variety of coherence methods have been developed over the years (Yao, Wan, and Xiao 2017), among which entity grid models is the most popular one proposed by (Barzilay and Lapata 2008). However, the discrete representation of entity grid suffer from the severe curse of dimensionality problem which limit its application on text summarization. (Nguyen and Joty 2017) presented a local coherence model based on a convolutional neural network that operates over the distributed representation of entity grid. However (Nguyen and Joty 2017) still rely on the name entity tools to identify the name entity of sentences first. (Li and Hovy 2014) use the recurrent or recursive neural network to obtain the distributed representation of sentences and then use pairwise ranking method to train the coherence model. (Li and Hovy 2014) select the window of sentences from original articles as positive examples, and random replacement of some sentences in the window as negatives examples. This model does not need any feature engineering, but this model is hard to capture the local entity transition for lack of cross-sentence local interaction. Our neural coherence model can be trained almost from scratch in an end-to-end fashion. It can also model the local entity transition and syntactical relation between sentences via different level cross-sentence local interaction.

## Neural Extractive Summarization Model

The prerequisite of using proposed reinforcement learning approach is to construct an neural extractive summarization model (NES). In this work, we build the state-of-the-art extractive summarization model based on SummaRuNNer proposed by (Nallapati, Zhai, and Zhou 2017). The extractive summarization model reads the document and selects a set of sentences sequentially as summary. Our contribution mainly lies in using the reinforcement learning to force the pretrained NES to produces more coherent and informative summary, which will be described in next section. This section mainly describes the architecture of the pretrained NES.

Given a document  $X = (x_1, x_2, \dots, x_n)$  that consist of  $n$  sentences, the neural extractive summarization models aim to output a sequence of binary decisions  $Y = (y_1, y_2, \dots, y_n)$ , where  $n$  denotes the number of sentences in the document and  $y_i \in \{0, 1\}$  indicates whether sentence  $x_i$  is selected. The  $i$ -th sentence is selected as part of summary if  $y_i = 1$ , and ignored otherwise.

The NES use a hierarchical neural network to encode the document. At the word-level, Convolutional neural network (CNN) is used to extract features of words and their context. Let  $x = (w_1, w_2, \dots, w_m)$  denotes a sentence with  $m$  words, and  $l$  denotes the size of word embedding. Then the sentence could be represented by a matrix  $\mathbf{M} \in \mathbb{R}^{m \times l}$ . Multiple convolution kernels with different kernel size are used to extract features.

$$\mathbf{f}_i^j = \mathbf{M}_{i:i+k_j-1} \odot \mathbf{W}_j + b_j,$$

where  $\odot$  represents element-wise product,  $\mathbf{W}_j, k_j, b_j$  denote the kernel weight matrix, kernel size and bias of the  $j$ -th convolution kernel respectively. The word  $w_i$  is represented by concatenating the feature maps  $\mathbf{f}_{w_i} = [\mathbf{f}_i^1; \mathbf{f}_i^2; \dots]$ . The representation of sentence  $x$  is represented by the mean of all its word features:

$$\mathbf{x} = \frac{1}{m} \sum_{i=1}^m \mathbf{f}_{w_i}.$$

At the sentence-level, we use a bi-directional gated recurrent unit (Bi-GRU) to model the context of sentences. Gated recurrent unit is a variant of recurrent neural network proposed by (Chung et al. 2014). It has two gates, an update gate  $\mathbf{z}_t$  and reset gate  $\mathbf{r}_t$ . The hidden state  $\mathbf{h}_t$  at time step  $t$  could be computed with following equations:

$$\mathbf{z}_t = \sigma(\mathbf{W}_z \mathbf{x}_t + \mathbf{V}_z \mathbf{h}_{t-1} + \mathbf{b}_z),$$

$$\mathbf{r}_t = \sigma(\mathbf{W}_r \mathbf{x}_t + \mathbf{V}_r \mathbf{h}_{t-1} + \mathbf{b}_r),$$

$$\hat{\mathbf{h}}_t = \tanh(\mathbf{W}_h \mathbf{x}_t + \mathbf{V}_h (\mathbf{r}_t \odot \mathbf{h}_{t-1}) + \mathbf{b}_h),$$

$$\mathbf{h}_t = (1 - \mathbf{z}_t) \odot \hat{\mathbf{h}}_t + \mathbf{z}_t \odot \mathbf{h}_{t-1},$$

where  $\mathbf{W}_z, \mathbf{W}_r, \mathbf{W}_h, \mathbf{V}_z, \mathbf{V}_r, \mathbf{V}_h$  and  $\mathbf{b}_z, \mathbf{b}_r, \mathbf{b}_h$  are parameters of GRU.

Using Bi-GRU, the representation of the  $t$ -th sentence  $\mathbf{x}_t$  is transformed to a forward hidden state  $\vec{\mathbf{h}}_t$  and a backward hidden state  $\overleftarrow{\mathbf{h}}_t$ . Both states are concatenated to form the contextual representation of  $t$ -th sentence  $\overleftrightarrow{\mathbf{h}}_t = [\vec{\mathbf{h}}_t; \overleftarrow{\mathbf{h}}_t]$ .

The entire document is represented as  $d$  by a non-linear transformation of the mean-pooling operation:

$$\mathbf{d} = \tanh(\mathbf{W}_d (\frac{1}{n} \sum_{t=1}^n \overleftrightarrow{\mathbf{h}}_t) + \mathbf{b}_d),$$

where  $\mathbf{W}_d$  and  $\mathbf{b}_d$  are parameters of the transformation, and  $n$  is the number of sentences in the document.

The probability of extraction decisions  $Y$  conditioned on document  $X$  could be factorized as  $\Pr(Y|X) = \prod_{t=1}^n \Pr(y_t|X, y_{1:t-1})$ . The probability of extracting the  $t$ -th sentence is given by the following equation:

$$\Pr(y_t = 1|X, y_{1:t-1}) = \text{MLP}(\overleftrightarrow{\mathbf{h}}_t, \mathbf{g}_{t-1}, \mathbf{d}), \quad (1)$$

where  $\mathbf{g}_t$  represents all sentences selected in history, and  $\text{MLP}(\cdot)$  means a multilayer perceptron that outputs a probability:

$$\text{MLP}(\overleftrightarrow{\mathbf{h}}_t, \mathbf{g}_{t-1}, \mathbf{d}) = \sigma(\tanh(\mathbf{W}_2 \tanh(\mathbf{W}_1 [\overleftrightarrow{\mathbf{h}}_t, \mathbf{g}_{t-1}, \mathbf{d}] + \mathbf{b}_1) + \mathbf{b}_2)),$$

where  $\mathbf{W}_1, \mathbf{W}_2, \mathbf{b}_1, \mathbf{b}_2$  are parameters of MLP and  $\sigma(\cdot)$  is the sigmoid function.

Since NES is trained with supervised learning, ground truth extraction labels  $(\hat{y}_1, \dots, \hat{y}_n)$  are available during training. Then the representation of previously extracted sentences at time  $t$  is:

$$\mathbf{g}_t = \mathbf{g}_{t-1} + \hat{y}_t \tanh(\mathbf{W}_g \overleftrightarrow{\mathbf{h}}_t). \quad (2)$$

The NES model is pretrained by minimizing the negative log-likelihood of the ground truth extraction labels

$$L_{\text{pretrain}}(\Theta) = - \sum_{i=1}^N \sum_{t=1}^{N_d} \hat{y}_t^i \log \Pr(y_t^i = 1|X_i, \hat{y}_{1:t-1}^i) + (1 - \hat{y}_t^i) \log \Pr(y_t^i = 0|X_i, \hat{y}_{1:t-1}^i).$$

## Reinforced Neural Extractive Summarization Model

In order to obtain the RNES model, the NES need to be trained using supervised learning strategy beforehand. And then NSE is forced to extract coherent and informative summary using reinforcement learning with the neural coherence and rouge as reward. In this section we first describe the REINFORCE algorithm, and then the neural coherence reward, ROUGE score reward.

### Reinforcement Learning

The problem of extractive summarization could be formulated as a reinforcement learning problem. The neural sentence extractor can be considered as an *agent* that extracts sentences sequentially from documents. At each time step  $t$ , the agent is in *state*  $s_t = (X, y_{1:t-1})$  which includes the document and previous selections. Agent would take a *action*  $y_t \in \{0, 1\}$  that decides whether sentence  $x_t$  is extracted or not. After the agent takes the action  $y_t$ , it may receive an immediate reward  $r_t$  that shows how good the action is. The reward could be delayed. When the agent finishes extracting

sentences from the whole document, it also receives an final reward  $r_{-1}$  that indicates performance of the entire action sequence  $(y_1, y_2, \dots, y_n)$ .

We use the REINFORCE algorithm to train our RNES model. It is a kind of policy gradient method proposed by (Williams 1992), and it maximizes agent’s performance by updating its policy. The policy is defined as the probability of taking an action at time  $t$  given the state, which is parameterized by  $\Theta$ :

$$\pi(a|s_t, \Theta) \triangleq \Pr(y_t = a|s_t, \Theta) \quad (3)$$

$$\triangleq \Pr(y_t = a|X, y_{1:t-1}, \Theta). \quad (4)$$

In this case,  $\Theta$  represents all the parameters in the neural sentence extractor. We use a shorthand  $\pi_\Theta$  to denote the policy  $\pi$  parameterized by  $\Theta$ . By applying Equation 1, we have:

$$\pi_\Theta(a = 1|s_t) = \text{MLP}_\Theta(\vec{\mathbf{h}}_t, \mathbf{g}_{t-1}, \mathbf{d}), \quad (5)$$

Let  $s_0 = X$  represents the initial state when no actions are taken yet, and  $v_{\pi_\Theta}(s)$  be the *value* function that represents the expected *return* starting with state  $s$  by following policy  $\pi_\Theta$ . Return at time  $t$  is defined as  $R_t = \sum_{i=t}^{\infty} \gamma^{i-t} r_i$ , where  $\gamma$  is the discount factor. The objective of REINFORCE is defined as maximizing the value of initial state  $v_{\pi_\Theta}(s_0)$ , or minimizing the loss  $L_{RF}(\Theta) = -v_{\pi_\Theta}(s_0)$ . Therefore, the parameters should be updated by the gradient of  $L_{RF}(\Theta)$  with respect to parameters  $\Theta$ :

$$\nabla L_{RF}(\Theta) = -\nabla v_{\pi_\Theta}(s_0) \quad (6)$$

$$= -\sum_{t=1}^n \gamma^t \Pr(s_t|s_0, \pi_\Theta) \sum_a q_{\pi_\Theta}(s_t, a) \nabla \pi_\Theta(a|s_t), \quad (7)$$

where  $q_{\pi_\Theta}(s, a)$  is the *action-value* function that represents the expected return after taking action  $a$  at state  $s$  with policy  $\pi_\Theta$ .

Since the state space is too large, it is infeasible to compute the exact value of the gradient. We use Monte Carlo sampling to approximate the gradient:

$$\nabla L_{RF}(\Theta) = -\mathbb{E}_{\tilde{s}_t \sim \pi_\Theta} \left[ \gamma^t \sum_a q_{\pi_\Theta}(\tilde{s}_t, a) \nabla \pi_\Theta(a|\tilde{s}_t) \right] \quad (8)$$

$$= -\mathbb{E}_{\tilde{y}_t, \tilde{s}_t \sim \pi_\Theta} \left[ \gamma^t q_{\pi_\Theta}(\tilde{s}_t, \tilde{y}_t) \frac{\nabla \pi_\Theta(\tilde{y}_t|\tilde{s}_t)}{\pi_\Theta(\tilde{y}_t|\tilde{s}_t)} \right] \quad (9)$$

$$= -\mathbb{E}_{\tilde{y}_t, \tilde{s}_t \sim \pi_\Theta} \left[ \gamma^t q_{\pi_\Theta}(\tilde{s}_t, \tilde{y}_t) \nabla \log \pi_\Theta(\tilde{y}_t|\tilde{s}_t) \right] \quad (10)$$

$$= -\mathbb{E}_{\tilde{y}_t, \tilde{s}_t \sim \pi_\Theta} \left[ \gamma^t \tilde{R}_t \nabla \log \pi_\Theta(\tilde{y}_t|\tilde{s}_t) \right], \quad (11)$$

where  $\tilde{s}_t$  and  $\tilde{y}_t$  are randomly sampled from  $\pi_\Theta$ ,  $\tilde{R}_t$  is the actual return received since  $\tilde{s}_t$  and  $\tilde{y}_t$ . Then parameters  $\Theta$  is updated as follows:

$$\Theta \leftarrow \Theta + \gamma^t \tilde{R}_t \nabla \log \pi_\Theta(\tilde{y}_t|\tilde{s}_t). \quad (12)$$

The definition of reward is crucial for reinforcement learning, because it determines the optimization direction.

**Algorithm 1** The overall training algorithm of RNES model.  $\alpha$  is the learning rate,  $\chi$  is a placeholder sentence for bootstrapping the coherence score of the first extracted sentence.

---

```

1:  $\Psi \leftarrow$  train the neural coherence model.
2:  $\Theta \leftarrow$  pretrain the neural sentence extractor with supervised learning.
3: loop
4:    $X, H \leftarrow$  sample a document-summary pair from corpus  $D$ 
5:    $\tilde{s}_0 \leftarrow X$ 
6:   Sample an episode  $\tilde{s}_1, \tilde{y}_1, \dots, \tilde{s}_n, \tilde{y}_n$  following  $\pi_\Theta$ 
7:    $previous \leftarrow \chi$  (a placeholder for empty start sentence)
8:   for each step  $t = 1 \dots n$  do
9:     if  $\tilde{y}_t = 1$  then
10:       $\tilde{r}_t \leftarrow \lambda \text{Coh}_\Psi(previous, x_t)$ 
11:       $previous \leftarrow x_t$ 
12:     else
13:       $\tilde{r}_t \leftarrow 0$ 
14:    $G \leftarrow$  the sequence of extracted sentences
15:    $\tilde{r}_{-1} = \text{ROUGE}(G, H)$ 
16:   for each step  $t = 1 \dots n$  do
17:      $\tilde{R}_t \leftarrow \sum_{i=t}^n \tilde{r}_i + \tilde{r}_{-1}$ 
18:    $\Theta \leftarrow \Theta + \alpha \tilde{R}_t \nabla \log \pi_\Theta(\tilde{y}_t|\tilde{s}_t)$ 

```

---

To ensure that neural extractive summarization model to extract coherent and informative summaries, the reward includes both coherence score and ROUGE score. The coherence score is exploited as immediate reward  $r_t$  and ROUGE as the final reward  $r_{-1}$ . Therefore, the algorithm is indeed maximizing the sum of both ROUGE and coherence score:

$$v_\pi(s_0) \triangleq \mathbb{E}_\pi[R_0|s_0] = \mathbb{E}_\pi[\tilde{r}_{-1} + \sum_{t=1}^n \tilde{r}_t] \quad (13)$$

$$= \mathbb{E}_\pi[\text{ROUGE}(\tilde{Y}) + \lambda \text{Coherence}(\tilde{Y})], \quad (14)$$

where  $\tilde{Y} = (\tilde{y}_1, \dots, \tilde{y}_n)$  is a sampled sequence of actions, and  $\tilde{r}$  is their corresponding rewards.  $\lambda$  is the coefficient that balances the two rewards. The neural coherence model for computing the coherence reward and Rouge score reward are introduced in the next subsection. Alg. 1 shows the REINFORCE algorithm to train our proposed RNES model.

### Neural Coherence Reward

We propose a neural coherence reward model which is based on convolutional neural network architecture. The neural coherence reward model has some advantages over traditional entity grid models. On one hand, the neural coherence model compute the coherence score almost from scratch. It only need take the original sentences as input. On other hand, The neural coherence reward model use the distributed text representation which can capture long range entity transitions. The architecture of the neural coherence reward model is shown in Figure . The Neural Coherence reward model is build on the ARC-II proposed by (Hu et al. 2014) for sentence matching.

Given two sentences  $S_X$  and  $S_Y$ , the neural coherence model can model the local interaction between two sentences. In Layer-1, it takes sliding windows on both sentences to model all the possible local coherence matching of two sentences. For segment  $i$  on  $S_X$  and segment  $j$  on




Figure 1: Architecture of neural coherence model which is build on ARC-II proposed by (Hu et al. 2014).

$S_Y$ , the local coherence matching of them is computed as Eq. 15

$$z_{i,j}^{(1)} \stackrel{\text{def}}{=} z_{i,j}^{(1)}(x, y) = g(\hat{\mathbf{z}}_{i,j}^{(0)}) \cdot \sigma(W^{(1)}\hat{\mathbf{z}}_{i,j}^{(0)} + b^{(1)}), \quad (15)$$

where  $W^1$  is the weight parameters for 1st layer and  $b^1$  is the bias.  $\sigma$  is the nonlinear function which we choose ReLU (Dahl, Sainath, and Hinton 2013).  $\hat{\mathbf{z}}_{i,j}^{(0)} \in \mathbb{R}^{2k_1 \times D_e}$  is got by concatenating the embeddings of words in  $S_X$  and  $S_Y$  sequentially:

$$\hat{\mathbf{z}}_{i,j}^{(0)} = [e(x_i)^\top, \dots, e(x_{i+k_1-1})^\top, e(y_j)^\top, \dots, e(y_{j+k_1-1})^\top]^\top. \quad (16)$$

$g(\hat{\mathbf{z}}_{i,j}^{(0)})$  is the gate function, It equals 0 if all the element of  $\hat{\mathbf{z}}_{i,j}^{(0)}$  equals 0 else 1.

Layer2 takes in the output of layer-1 and performs a max-pooling in each dimension of the cells in non-overlapping  $2 \times 2$  windows.

$$z_{i,j}^{(2)} = \max(\{z_{2i-1,2j-1}^{(1)}, z_{2i-1,2j}^{(1)}, z_{2i,2j-1}^{(1)}, z_{2i,2j}^{(1)}\}). \quad (17)$$

Following Layer-2, there is more convolution and max-pooling layers, analogous to that of convolutional architecture for image input (LeCun and Bengio 1995). Finally, we obtain the fixed length vector  $h$  with dimension  $m$  and feed  $h$  into multi-layer perceptron (MLP), a nonlinear function  $\tanh$ , to compute coherence score of the two sentences.

$$\text{Coh}(S_X, S_Y) = \tanh(Wh + b). \quad (18)$$

Where  $W \in \mathbb{R}^{1 \times m}$  is the weight parameters and  $b$  is the bias. Hence, the coherence model will output a coherence score  $\text{Coh}(S_X, S_Y) \in [-1, 1]$  for sentence pairs  $S_X, S_Y$ .

From the first layer, the neural coherence model can learn the local interaction of  $S_X$  and  $S_Y$ , so as to capture the local entities transition of two sentences. And also, it can obtain higher level and wider entities transition representation of  $S_X$  and  $S_Y$  with more convolution and max pooling layers.

For training the neural coherence model, we use a discriminative training strategy with a large margin objective. Suppose we are given the following triples  $(S_X, S_Y^+, S_Y^-)$ , we have the ranking-based loss as objective:

$$L_\Theta(S_X, S_Y^+, S_Y^-) = \max(0, m + \text{Coh}(S_X, S_Y^-) - \text{Coh}(S_X, S_Y^+)) \quad (19)$$

Where  $\Theta = m \leq 1$  controlling the margin in training. In the experiments, we use  $m = 1.0$ . The model is trained by minimizing the above objective, to encourage the model to assign higher coherence score for coherent sentence pairs  $(S_X, S_Y^+)$  than  $(S_X, S_Y^-)$ . We use stochastic gradient descent (SGD) to optimize the model parameter.

It is easy to construct large enough data to train neural coherence model. Given a document with sequence of  $n$  sentences  $S_1, S_2, \dots, S_n$ . We can obtain the following triples  $(S_l, S_{l+1}, S_j)$ , where  $0 < l < n$ .  $S_j$  is randomly sampled from document and  $j \notin [l, l+1]$ .

## ROUGE Score Reward

ROUGE score is used as the final reward to ensure that the neural extractor extracts reasonably informative sentences. Given a sequence of sampled decisions  $(\tilde{y}_1, \dots, \tilde{y}_n)$ , we can get the sequence of extracted sentences  $G$ . Since the dataset comes with news highlights written by human editors, these manual highlights  $H$  is treated as the reference summary. Then ROUGE score  $\text{ROUGE}(G, H)$  could be computed and used as a reward for the entire sampled decisions:

$$r_{-1}(\tilde{y}_1, \dots, \tilde{y}_n) = \text{ROUGE}(G, H).$$

Multiple variants of ROUGE score are proposed by (Lin 2004). Among them, ROUGE-1, ROUGE-2 and ROUGE-L are the most commonly used ones. ROUGE- $n$  recall between an extracted summary and a reference summary can be computed as follows:

$$\text{ROUGE-}n = \frac{\sum_{s \in \text{reference summary}} \sum_{\text{gram}_n \in s} \text{Count}_{\text{match}}(\text{gram}_n)}{\sum_{s \in \text{reference summary}} \sum_{\text{gram}_n \in s} \text{Count}(\text{gram}_n)},$$

where  $n$  stands for the length of  $n$ -gram,  $\text{Count}_{\text{match}}(\text{gram}_n)$  is the maximum number of  $n$ -grams co-occurring in both the extracted summary and the reference. Similarly we could compute the ROUGE- $n$  precision and F1. ROUGE-1 and ROUGE-2 are special cases of ROUGE- $n$  in which  $n = 1$  or  $n = 2$ . ROUGE-L is instead computed based on the length of longest common subsequence between the candidate summary and the reference. Since using only variant of ROUGE as reward for training RNES may not increase its performance on other ROUGE variants, we used a combination of ROUGE variants as reward:

$$\text{ROUGE}(G, H) = w_1 \text{ROUGE-1}(G, H) + w_2 \text{ROUGE-2}(G, H) + w_l \text{ROUGE-L}(G, H),$$

where weights  $w_1, w_2$  and  $w_l$  are hyper-parameters.

The overall training algorithm is illustrated in the Algorithm 1. Since REINFORCE algorithm converges very slowly, we pretrain the neural sentence extractor with supervised learning. The neural coherence model is also trained and then fixed for the coherence scoring. During the training

of REINFORCE, an sequence of actions and states are sampled according to the policy. Then the coherence model and ROUGE are used for computing the rewards. The parameters of neural sentence extractor  $\Theta$  is then updated according to Equation 12.

## Experiments and Results

We used the CNN/Daily Mail dataset ~~originally proposed by~~ <sup>introduced</sup> (Hermann et al. 2015) ~~for machine comprehension task~~. This dataset contains news documents and highlights crawled from CNN and Daily Mail website, and (Nallapati et al. 2016) ~~propose to use it for summarization task~~. <sup>and</sup> It is commonly used by works of <sup>SH</sup> extractive summarization (Cheng and Lapata 2016; Nallapati, Zhai, and Zhou 2017) and abstractive summarization (Nallapati et al. 2016; See, Liu, and Manning 2017). We used the scripts provided by (Hermann et al. 2015) to download the dataset. It contains 287,226 documents for training, 13,368 documents for validation and 11,490 documents for test. We use the NLTK package (Bird, Klein, and Loper 2009) to segment the sentences. Since the dataset only contains manual summaries and do not have extractive labels, a greedy algorithm similar to (Nallapati, Zhai, and Zhou 2017) <sup>is used to generate extraction labels for supervised learning</sup> <sup>training of NES</sup>.

## Results of the Neural Coherence Model

The coherence model need to be trained <sup>before it could be used</sup> in order to use the coherence score as reward in the REINFORCE algorithm. In <sup>SH</sup> experiment, we uses 64-dimensional word embeddings which is randomly initialized and updated in the process of supervised training, and the sizes of all its convolutional kernels are set to be 3. In the first convolution layer, it has 128 filters. It is followed by two layers of convolution, with 512 and 256 filters respectively. <sup>Each of the convolutional layers</sup> ~~is followed by a max-pooling layer with size 2 and stride 2.~~ <sup>which are</sup> The final two fully-connected layers have 512 and 256 hidden units respectively. The maximum sentence length is set 50, and sentences longer than this length would be truncated, less will be padded with 0s. The coherence model is trained with stochastic gradient descent (SGD) with batch size 64 and learning rate 0.1.

The training triplets are sampled from the CNN/Daily Mail dataset. The  $S_X$  and  $S_Y^-$  are adjacent sentences sampled from the documents, and  $S_Y^+$  is a sentence randomly sampled such that  $S_Y^- \neq S_Y^+$ . To make the task more difficult so that coherence model finds more fine grained coherence patterns,  $S_Y^-$  is sampled from the same document as  $(S_X, S_Y^+)$  and its distance is less than 9 sentences from  $S_Y^+$ .

The model is tested on 23K positive pairs sampled from the test set, each accompanied with four negative samples. If the model gives a higher score to the positive sample than the negative sample, it is considered correct. The accuracy is <sup>70%</sup> ~~70%~~ versus 50% accuracy for random guess, which indicates the neural coherence model can capture the cross-sentence coherence effectively.

We also conducted empirical studies about the output of coherence model. Table 1 shows some examples of coherence scoring. The first example shows that the model can

Table 1: Example outputs of neural coherence model.

Sentences	Score
$S_X$ : Terry's career as a <b>photographer</b> came after he failed to make it as a punk rock musician. $S_Y^-$ : <b>He</b> got his first big break in 1994 with a <b>shoot</b> for Vibe magazine. $S_Y^+$ : The photographer has also directly music videos in his time.	<b>0.9885</b> 0.5198
$S_X$ : These days we are increasingly using outdoor space for the occasional barbecue or to relax in a hot tub rather than for tending <b>flowers</b> , according to researchers. $S_Y^-$ : <b>As a result</b> , only a handful of traditional <b>flowers</b> still grow in English country gardens, with the average one usually containing a mere four species - <b>daffodils, crocuses, roses and tulips</b> . $S_Y^+$ : Sir Roy Strong, the landscape designer and former director of the Victoria and Albert Museum, told the Sunday Times: 'British people used to take pride in having neat gardens with lots of flowers.'	<b>0.8934</b> -0.0067
$S_X$ : The same <b>survey</b> recently showed that university pupils in Britain have an average of 8.2 sexual partners by the time they reach the middle of their higher education. $S_Y^-$ : A new survey of university students has revealed that they have had an average of 8.2 sexual partners (picture posed by models) $S_Y^+$ : <b>According to the answers they received</b> , 22 per cent of students didn't lose their virginity until they were 18 years old, with the second most popular age to have sex for the first time being 16.	0.0021 <b>0.9999</b>
$S_X$ : Christie has also forcefully denounced marijuana legalization, saying last March, ' <b>I don't favor</b> legalization. $S_Y^-$ : <b>I don't favor</b> recreational use. $S_Y^+$ : <b>And I don't favor</b> the use of marijuana as a medicine.' $S_Y^-$ : The governor is on a four-day swing through the first in the nation primary state as he explores a run for the Republican nomination for president.	<b>0.9995</b> <b>0.9996</b> -0.0344

exploit co-reference for coherence modeling. The model can also find correlations between semantically related words such as "photographer" and "shoot" (example 1), "survey" and "answers" (example 3). Furthermore, the coherence model is able to discover patterns such as "As a result ...", "According to ..." and parallel construction "I don't favor ...", which are useful for coherence estimation. The third example also shows that, adjacent sentences are not always coherent, because the documents may contain image or video captions that are embedded in the article. <sup>The model gives low score to the  $S_Y^-$  in example 3, which is reasonable because sentences containing same information is less likely to be adjacent.</sup> <sup>adjacent vs. coherent</sup>

## Results of Reinforced Neural Extractive Model

For the NES model, we use 128-dimensional word embeddings and the vocabulary size is 150,000. The convolution



kernels have size 3, 5, 7 with 128, 256, 256 filters respectively. We set the hidden state size of sentence-level GRU to 256, and the document representation size to 512. The MLP has two layers, with 512 and 256 hidden units respectively. We fix the maximum sentence length to 50 and the maximum number of sentences in a document to 80. Sentences or documents that are longer than the maximum length are truncated to fit the length requirement.

The model is trained with stochastic gradient descent (SGD) with batch size 64. At test time, the NES model selects sentences using beam search with beam size 10. When doing supervised training of the NES, ground truth extraction labels are used for computing the cross-entropy loss. The labels are generated from the dataset by greedily selects sentences to maximize the ROUGE score with respect to manual highlights.

During the training of RNES using reinforcement learning, both the neural coherence model and ROUGE scorer are used to compute the reward. As shown in Equation 14, the hyper-parameter  $\lambda$  is used to balance between the two objectives. In our experiments, we tried  $\lambda = 1.0, 0.1, 0.01, 0.05$ . It is found that when  $\lambda > 0.01$ , the model favor coherence too much that ROUGE degrades rapidly and the model eventually converges to a policy that selects consecutive sentences that are not informative. However, when  $\lambda \leq 0.01$ , the ROUGE objective overpowers coherence, and the coherence rewards drop to approximately zero. We found that  $\lambda = 0.01$  is a good trade-off such that both rewards increase and eventually converge.

To compare with previous works, we adopt the same evaluation metrics as in (Nallapati, Zhai, and Zhou 2017). We use full length F1 of ROUGE-1, ROUGE-2 and ROUGE-L to evaluate our model. Table 2 shows the performance comparison between our models and current state-of-the-art. Both of our RNES model (with coherence and without coherence as reward) outperform current state of the art model by a large margin. This indicates that our proposed RNES is able to extract better summaries than previous works.

RNES without the neural coherence reward achieves slightly higher ROUGE score than the one trained with neural coherence as reward. Since ROUGE is simply computed based on n-gram and is ignorant of the coherence between sentences, enhancing coherence may lead to a drop of ROUGE. The result indicates that our model finds a good balance between the two objectives, while still outperforming previous works in terms of ROUGE.

Table 2: Performance comparison on CNN/Daily Mail test set, evaluated with full length F1 ROUGE scores (%).

Model	R-1	R-2	R-L
Lead-3	39.2	15.7	35.5
(Nallapati et al. 2016)	35.4	13.3	32.6
(Nallapati et al. 2017)	39.6	16.2	35.3
(See et al. 2017)	39.53	17.28	35.38
RNES w/o coherence	<b>41.25</b>	<b>18.87</b>	<b>37.75</b>
RNES w/ coherence	40.95	18.63	37.41

To further understand the performance of our models, we

also conduct qualitative evaluation. We sample 50 documents from the test set and ask three people to evaluate the summaries extracted by our systems trained with or without coherence reinforcement. The judges are asked to rank the outputs of two models, giving rank 1 to the better one and rank 2 to the other. The evaluation considers both informativeness and coherence aspect of the summary, and an overall ranking will also be provided by the judges. Table 3 compares the result of human evaluation. It shows that our model trained with coherence is ranked lower than the one trained without coherence in all aspects. The result indicates that, the introduction of coherence reinforcement effectively improves the coherence of summaries, and therefore improves the overall quality. It is surprising that the informativeness of our coherence-enhanced model is also slightly improved.

Table 3: Comparison of human evaluation in terms of informativeness, coherence and overall ranking. Lower is better.

	informativeness	coherence	overall
RNES w/o coherence	1.183	1.325	1.492
RNES w/ coherence	<b>1.125</b>	<b>1.092</b>	<b>1.209</b>

## Conclusion

In this paper, we proposed a Reinforced Neural Extractive Summarization model to extract coherent and informative summary for single document. We also proposed a neural coherence model to capture the cross sentence coherence. Empirical results show that the proposed RNES model can balance between the cross-sentence coherence and informative importance of sentence effectively. Compared to state of the art models, The RNES model can output higher quality summary. In the future work, we will focus on improving the performance of our neural coherence model and introduce human knowledge into the RNES.

Table 4: Examples of extracted summary.

<i>Title:</i> Earth's newest supercontinent is taking shape: Land masses are already drifting together to form 'Amasia'
<i>Reference:</i> Peter Spinks from the Sydney Morning Herald reported on Amasia. Within 200 million years, he said the new supercontinent will form. One researcher recently travelled to Nepal to gather further information. He spotted that India, Eurasia and other plates are slowly moving together.
<i>Our model:</i> That's according to one researcher who travelled to the country to study how the Indian and Eurasian plates are moving together. And using new techniques, researchers can now start examining the changes due to take place over the next tens of millions of years like never before. Earth's continents are slowly moving together (left), and in 50 to 200 million years they are expected to form a new supercontinent called Amasia (right). In 2012 a study suggested this may be centered on the North Pole. The idea that Earth is set to form a new supercontinent-dubbed Amasia - is not new.
<i>Our model+coherence:</i> The earthquake disaster in Nepal has highlighted how Earth's land masses are already in the process of forming a <b>new supercontinent</b> . That's according to one researcher who travelled to the country to study how the <b>Indian and Eurasian plates are moving together</b> . And using new techniques, researchers can now start examining the changes due to take place over the next tens of millions of years like never before. Earth's continents are slowly moving together (left), and in 50 to <b>200 million years</b> they are expected to form a new supercontinent called Amasia (right).
<i>Title:</i> Give us 500 or we'll chop up your dog: Grandmother 'sick and scared' after chilling ransom demand for her stolen Westie Ben
<i>Reference:</i> Vivienne Garton, 65, was devastated when her 10-year-old dog was stolen. She started appeal to find him but received a call from anonymous man. Caller threatened to cut up her West Highland terrier unless she paid 500. He gave address of empty house and police are investigating the blackmail.
<i>Our model:</i> Vivienne Garton, 65, was devastated her 10-year-old West Highland terrier Ben was snatched from her garden in Knowle West, Bristol, earlier this month. She began an appeal to find him and put up posters near her home, but shortly after she received two phone calls from an anonymous man threatening to cut up her pet unless she paid 500. <i>Scared: Vivienne Garton, 65, has been left devastated after her West Highland terrier Ben was stolen.</i>
<i>Our model+coherence:</i> A grandmother was left feeling 'sick and scared' after she received a chilling ransom demand for her stolen dog. Vivienne Garton, 65, was devastated her 10-year-old West Highland terrier Ben was snatched from her garden in Knowle West, Bristol, earlier this month. She began an appeal to find him and put up posters near her home, but shortly after she received two phone calls from an anonymous man threatening to cut up her pet unless she paid 500.



## References

- Ayana; Shen, S.; Liu, Z.; and Sun, M. 2016. Neural headline generation with minimum risk training. *CoRR* abs/1604.01904.
- Bahdanau, D.; Cho, K.; and Bengio, Y. 2014. Neural machine translation by jointly learning to align and translate. *CoRR* abs/1409.0473.
- Barzilay, R., and Lapata, M. 2008. Modeling local coherence: An entity-based approach. *Comput. Linguist.* 34(1):1–34.
- Berg-Kirkpatrick, T.; Gillick, D.; and Klein, D. 2011. Jointly learning to extract and compress. In *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies - Volume 1*, HLT '11, 481–490. Stroudsburg, PA, USA: Association for Computational Linguistics.
- Bird, S.; Klein, E.; and Loper, E. 2009. *Natural Language Processing with Python*. O'Reilly Media, Inc., 1st edition.
- Cheng, J., and Lapata, M. 2016. Neural summarization by extracting sentences and words. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics, ACL 2016, August 7-12, 2016, Berlin, Germany, Volume 1: Long Papers*.
- Chung, J.; Gulcehre, C.; Cho, K.; and Bengio, Y. 2014. Empirical evaluation of gated recurrent neural networks on sequence modeling. *arXiv preprint arXiv:1412.3555*.
- Dahl, G. E.; Sainath, T. N.; and Hinton, G. E. 2013. Improving deep neural networks for lvsr using rectified linear units and dropout. In *Proceedings of ICASSP*.
- Erkan, G., and Radev, D. R. 2004. Lexrank: Graph-based lexical centrality as salience in text summarization. *J. Artif. Int. Res.* 22(1):457–479.
- Hermann, K. M.; Kocisky, T.; Grefenstette, E.; Espeholt, L.; Kay, W.; Suleyman, M.; and Blunsom, P. 2015. Teaching machines to read and comprehend. In *Advances in Neural Information Processing Systems*, 1693–1701.
- Hu, B.; Lu, Z.; Li, H.; and Chen, Q. 2014. Convolutional neural network architectures for matching natural language sentences. In *Advances in Neural Information Processing Systems 27*, 2042–2050.
- Hu, B.; Tang, B.; Chen, Q.; and Kang, L. 2016. A novel word embedding learning model using the dissociation between nouns and verbs. *Neurocomput.* 171(C):1108–1117.
- Hu, B.; Chen, Q.; and Zhu, F. 2015. LCSTS: A large scale chinese short text summarization dataset. *CoRR* abs/1506.05865.
- LeCun, Y., and Bengio, Y. 1995. Convolutional networks for images, speech and time series. *The Handbook of Brain Theory and Neural Networks* 3361.
- Li, J., and Hovy, E. H. 2014. A model of coherence based on distributed sentence representation. In Moschitti, A.; Pang, B.; and Daelemans, W., eds., *EMNLP*, 2039–2048. ACL.
- Lin, C.-Y. 2004. Rouge: A package for automatic evaluation of summaries. In Marie-Francine Moens, S. S., ed., *Text Summarization Branches Out: Proceedings of the ACL-04 Workshop*, 74–81. Barcelona, Spain: Association for Computational Linguistics.
- Mikolov, T.; Chen, K.; Corrado, G.; and Dean, J. 2013. Efficient estimation of word representations in vector space. *CoRR* abs/1301.3781.
- Nallapati, R.; Zhou, B.; Gulcehre, C.; and Xiang, B. 2016. Abstractive Text Summarization using Sequence-to-sequence RNNs and Beyond. In *Proceedings of the 20th {SIGNLL} Conference on Computational Natural Language Learning*.
- Nallapati, R.; Zhai, F.; and Zhou, B. 2017. Summarunner: A recurrent neural network based sequence model for extractive summarization of documents. In *Proceedings of the Thirty-First AAAI Conference on Artificial Intelligence, February 4-9, 2017, San Francisco, California, USA.*, 3075–3081.
- Nayeem, M. T., and Chali, Y. 2017. Extract with order for coherent multi-document summarization. *TextGraphs-11* 51.
- Nguyen, D. T., and Joty, S. R. 2017. A neural local coherence model. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics, ACL 2017, Vancouver, Canada, July 30 - August 4, Volume 1: Long Papers*, 1320–1330.
- Nguyen, K.; Boyd-Graber, J.; and Daume, H. I. 2017. Reinforcement learning for bandit neural machine translation with simulated human feedback. In *Empirical Methods in Natural Language Processing*.
- Paulus, R.; Xiong, C.; and Socher, R. 2017. A deep reinforced model for abstractive summarization. *CoRR* abs/1705.04304.
- Ranzato, M.; Chopra, S.; Auli, M.; and Zaremba, W. 2015. Sequence level training with recurrent neural networks. *CoRR* abs/1511.06732.
- Rush, A. M.; Chopra, S.; and Weston, J. 2015. A neural attention model for abstractive sentence summarization. In *EMNLP*, 379–389.
- See, A.; Liu, P. J.; and Manning, C. D. 2017. Get To The Point: Summarization with Pointer-Generator Networks. *arXiv:1704.04368 [cs]*. arXiv: 1704.04368.
- Tang, D.; Qin, B.; and Liu, T. 2015. Document modeling with gated recurrent neural network for sentiment classification. In *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*, 1422–1432. Lisbon, Portugal: Association for Computational Linguistics.
- Williams, R. J. 1992. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine learning* 8(3-4):229–256.
- Yao, J.; Wan, X.; and Xiao, J. 2017. Recent advances in document summarization. *Knowledge and Information Systems* 1–40.